

EdgePilot AI - Project Record

Every fix, its owner, the structural comparison, and the architecture payoffs

EdgePilot AI - Team 15
Documentation set, 23 August 2026
1 document in this volume

EdgePilot AI - Project Record

Every change made to this codebase, the part of the system it touched, what was wrong before it, and what the structure looks like now against the version the team shipped.

Project: EdgePilot AI - compare local and cloud AI deployment with real measurements. **Team:** 15. **This working copy:** Documents/Air On Accountant Backup/EdgePilot testing/app. **Compared against:** Desktop/Air on Accountants/edgepilot-ai-main (1)/edgepilot-ai-main, the team's repository after their updates. **Record written:** 23 August 2026, at commit b13af6e.

1. How to read this document

Section 3 is the fix list. Section 4 is the structural comparison. Section 5 is the argument for why this structure is the better one, and section 6 is the honest list of what it still does not do.

How a change is placed. Every fix below names the **work package** it lands in, not a person. The purpose is navigational: it tells you which part of the system to open, and therefore which conversation the change belongs to. The areas are the ones the repository already draws - the module a file sits in - so nothing here depends on anyone's recollection. Where a fix crosses two areas, both are named.

Verification method. Every "before" state in section 3 is a real file in the comparison copy on the Desktop, and every "after" state is a real file here. A claim of the form "this returned a constant" can be checked by opening both. Line counts come from diff between the two trees.

2. Where each work package lives

Use this to find the part of the system a change belongs to, and therefore which conversation it needs.

Benchmark and provider layer - measurement, the provider adapters, scoring and evidence. src/modules/benchmark/, the benchmarks, providers, readiness, comparisons and local-runtime endpoints under src/app/api/v1/, scripts/benchmark/, scripts/evaluation/, docs/benchmark/, docs/local-model-setup.md.

Integration and API surface - routes, shared types, the database schema and the internal documentation. src/app/api/v1/workloads/, src/shared/types/, prisma/, docs/internal/.

Product UI and benchmark dashboard - the four-step user journey and its styling. src/app/dashboard/, src/components/dashboard/, docs/dashboard/.

Vision benchmark - image classification end to end, including the dataset. src/modules/vision-benchmark/, src/app/vision-benchmark/, src/components/vision/, scripts/vision-benchmark/, docs/vision-benchmark/, datasets/.

Device knowledge and quality - removed in this pass; see change 22. The former src/modules/device/, src/app/api/v1/devices/, device-profile.json.

DevOps - the container, CI, the shared-database guards and the evidence page. Dockerfile, .dockerignore, .github/workflows/, scripts/db/, src/app/evidence/.

These are the divisions the repository itself already uses: AI_USAGE.md is organised one section per work package, and docs/benchmark/README.md and docs/dashboard/README.md are each written as a module guide.

3. The fix list

Forty-one changes, grouped by what they were about. Small ones are included - a one-character regex fix is in here, because it was the difference between a feature working and a feature that had never once worked.

3.1 Measurement correctness

These are the changes that decide whether a number in this application means anything.

1. hardwareFit was a constant. *Area: benchmark layer.* Every device on earth received `hardwareFit: 50`, from a Raspberry Pi to a threadripper. It is visible in the committed evidence: `evidence/benchmark/measured-ollama-llama3.2_latest.json` still records `"hardwareFit": 50`, because that artefact predates the fix. Replaced with `src/modules/benchmark/core/services/HardwareFitAssessor.ts` (229 lines, new), which scores *placement* from what the runtime reports, and `src/modules/benchmark/infrastructure/OllamaResidencyProbe.ts` (107 lines, new), which reads it. Ollama's `/api/ps` gives `size` (total resident bytes) and `size_vram` (how many of those are on the GPU); `size_vram / size` is the fit signal, and it is reported by the runtime rather than inferred by us.

2. Hardware fit deliberately does not score speed. *Area: benchmark layer.* Latency is already measured and already a readiness component. If hardware fit also punished CPU execution, a slow run would be penalised twice and the total would stop meaning anything. Recorded in the file's own header so the next person does not "fix" it back.

3. VRAM taken and spill share are now stated in words. *Area: benchmark layer.* The assessment says how much VRAM the model took and, on a partial offload, how many bytes did not fit and what percentage of the model that was. Where the runtime does not report card capacity, the assessment says so rather than presenting a share of an unknown total.

4. Cold start contaminated every first iteration. *Area: benchmark layer.* A model that is not resident has to be loaded from disk before it can answer. Measured on the same 8B model in the same session: time-to-first-token was **36,445 ms cold** and **369 ms warm** - a factor of 99 - and readiness swung from 75 to 90 on load state alone. Averaging that first iteration into a five-iteration mean does not describe the model; it describes whether someone had used it recently. Fixed in `BenchmarkRunner.ts`: the runner now runs N+1 iterations and discards the first. The user is told their chosen iteration count will be raised by one, the discarded cold boot is reported with its own duration and its own row in the results, and the database records it as `warmup: true` rather than deleting it.

5. The success chain counted the discarded iteration. *Area: benchmark layer.* After change 4, a request for three iterations reported "4/4 iterations succeeded". The measured set is now `responses.slice(1)` and success, failure-code dominance and the summary all read from that. The fix went into the runner rather than the test that caught it.

6. The residency probe result was used without a null guard. *Area: benchmark layer.* The catch produced the correct value for the wrong reason, which would have masked any other fault in the probe. Now before `=== null ? null : before.residentBytes !== null && before.residentBytes > 0`.

7. Privacy was averaged into the readiness score. *Area: benchmark layer.* A coarse label was mapped to 30/60/100 and blended with latency and reliability. That is a category error twice over: it implies a vendor is "twice as private" as another, and it lets a fast provider average away a data policy that should have disqualified it outright. `ReadinessCalculator.ts` dropped from a five-term average to four measured performance terms, and privacy became an ordinal class produced by `PrivacyAssessor.ts` (314 lines, new) that annotates or disqualifies a recommendation instead of nudging a number.

8. A mean was being reported as a result. *Area: benchmark layer.* `ComparisonReport.ts` (506 lines, new) checks every dimension against the per-iteration samples behind it. If the entrants' observed ranges overlap, no

winner is declared. At five iterations a 10% gap in means routinely sits inside a 3x spread. A worked consequence: after cold starts were removed, latency did **not** become a cleanly separated dimension - separability *fell* from 0.82 to 0.58, because removing the cold start shrank the gap more than it shrank the spread. The report says so rather than declaring a winner.

9. Two local models were about to be compared in parallel. *Area: benchmark layer.* For two cloud providers, parallel is correct. For two local models it destroys the measurement: both want the same GPU, the runtime may evict one to make room for the other, and the latency that comes back describes contention. `ComparisonPlanner.ts` (141 lines, new) decides execution mode before anything runs, and `RunComparison.ts` honours the plan instead of always calling `Promise.all`.

10. Nothing stopped a text model being compared to an embedding model. *Area: benchmark layer.* `ModelModality.ts` (167 lines, new) classifies from Ollama's `families` array. Vision (`clip`, `mlama`, `llava`, `qwen2vl`, `vision`) and embedding (`bert`, `nomic-bert`, `gte`, `jina-bert`) are identified positively; text is the residual and is always marked `inferred`, because "no vision family present" is not proof of anything.

11. task_type was a column nothing read. *Area: benchmark layer, over the integration layer's workload schema.* You could register an `image_recognition` workload and benchmark an embedding model against it and no part of the system objected. `TaskCompatibility.ts` (174 lines, new) turns the declaration into a constraint. The rule is deliberately asymmetric: a vision model can do text, a text model cannot do vision, an embedding model can do neither.

12. Timeouts were too low for a cold local model, and for cloud. *Area: benchmark layer.* Raised for the local path so a genuine cold load is not recorded as a timeout, and raised for the online path in the same pass.

13. code_generation no longer implies the code works. *Area: benchmark layer.* `shapeOf('code_generation')` returns text that says so: nothing in this system compiles or runs the output, so "it produced code" is not "the code works".

3.2 Identity, ownership and privacy

14. Users were being asked to paste UUIDs by hand. *Area: dashboard, and integration.* The setup step had a field for a workload UUID because the endpoint that should have created one did not persist. Replaced by `src/components/dashboard/session.ts` (92 lines, new): the browser generates a durable id, stores it in `localStorage`, sends it as `x-edgepilot-session` on every request, and **never renders it**. A user should no more be asked to see their own row-ownership key than to type their own cookie.

15. POST /api/v1/workloads never called the database. *Area: integration.* The committed version returned the literal string `temp-workload-id`, so nothing the dashboard saved reached Postgres and every benchmark run afterwards failed with 404 `Workload not found`. The route is now 126 lines against the original 53 and writes real rows. The original is kept beside it as `route.ts.scaffold.bak` and is referenced by name in the new file's header.

16. Every row was attributed to one shared user. *Area: integration.* Two routes attributed everything to `local-sandbox@edgepilot.invalid`. Harmless on one laptop; a data leak the moment it is hosted, because GET would hand every visitor every other visitor's workloads. `src/lib/sessionOwner.ts` (76 lines, new) upserts a User per session id. It fails closed: a request without the header gets 400, never a fallback to a shared user.

17. Session activity had nowhere to be recorded, and no redaction rule. *Area: benchmark layer.* `src/core/logging/SessionLog.ts` (407 lines, new) records what the application did, exportable as JSON, with no account attached. Prompt *content* is never recorded by default - only a length and a digest. A log that wrote out the same prompt the egress warning had just flagged would have defeated its own warning. Logging is opt-in by header: a request with no session id is not logged anonymously, it is not logged at all.

18. The privacy report arrived after the prompt had already left. *Area: benchmark layer.* `egress-warning.ts` (133 lines, new) answers the question before the run: is this prompt about to leave the machine, and to whom. It fails closed on a null endpoint.

19. The egress warning object contained the prompt text. *Area: benchmark layer.* Found by a test written in the same session. An object whose job is to warn that text is about to leave the machine must not carry a copy of that text. It now reports the character count and nothing else. An older assertion in `privacy-assessor.test.ts` had pinned the opposite behaviour and was reversed, with the reasoning recorded at the call site.

20. The privacy catalogue ships empty on purpose. *Area: benchmark layer.* Every value in `privacy-catalogue.ts` starts `null` and `unverified`. These are policy facts from terms of service, they differ between free and paid tiers, and they change without notice. Filling them in from memory is how an evidence-based tool starts publishing stale claims confidently. Filling them in is a human task and is listed in section 6.

21. The login page was removed. *Area: integration.* There are no accounts, so an auth surface was a door onto an empty room. `src/app/auth/` and `src/app/api/auth/` are gone.

3.3 Removing the device module

22. The device profile was a client-side self-report that nothing could check. *Area: device.* The user typed how much RAM they had; the application stored it; no measurement ever used it. Once hardware fit became `size_vram / size` read from the runtime, the typed value had no consumer and could only disagree with reality. Removed to the roots rather than left nullable: the UI panel, the API route, the module, the types, the tests, and the database.

23. The database migration. *Area: device; the migration came out of the benchmark and DevOps work.* `prisma/migrations/20260823035500_drop_devices/` drops two foreign keys, the `benchmarks_device_id_idx` index, the `benchmarks.device_id` column and the `devices` table. Tested on the local sandbox database before it was applied to the shared Neon instance, and applied there with `migrate deploy - never migrate dev`, which can offer to reset a shared database.

24. Two dev scripts still referenced devices, and next build caught them. *Area: benchmark layer.* `scripts/dev/show-last-run.ts` read `run.deviceId` and `scripts/dev/seed-demo-rows.ts` called `prisma.device.create`. Jest never touched either file. This is the reason `npm run build` is in the verification list and not just `npm test`.

3.4 The vision benchmark

25. The live evidence path had never worked, once, ever. *Area: vision.* `evidence-store.ts` validated filenames against `/^[a-z0-9._-]+\.json$/ - lowercase only -` and then built the filename from an ISO timestamp containing an uppercase T and Z. Every live run failed to write its evidence. One `.toLowerCase()` fixed it. It is the smallest change in this document and it is the reason any live measurement exists at all.

26. The first vision sample paid the cold-start cost. *Area: vision; the pattern came from the benchmark runner.* The executor now runs a warm-up classification before the measured set. Median latency across the dataset fell from 1322.1 ms to 529.6 ms between the two live runs - the same model, the same images, the difference being that sample one was no longer carrying the model load.

27. The warm-up was wrong three times before it was right. Recorded because the mistakes are instructive. It consumed a deterministic test clock (fixed by calling `provider.classify()` directly rather than the timed wrapper); it fired on cloud providers, which meant a billable call and one extra image sent to a third party (fixed: local providers only, plus an explicit `warmUp?: boolean`); and it ate scripted responses in two existing tests (fixed by opting those two out with a comment, and adding two new tests that cover the behaviour properly).

28. Users can now bring their own dataset, and it never leaves their machine. *Area: vision.* `src/components/vision/DatasetUpload.tsx` (403 lines, new) takes a folder-per-class upload. `browser-image-processor.ts` (181 lines, new) resizes and hashes in the tab; `browser-ollama-provider.ts` (171 lines, new) talks from the tab straight to the user's own Ollama on localhost. No uploaded image ever reaches the EdgePilot server. That is not a convenience decision - benchmarking confidential data is this project's stated non-goal, and shipping images to a server to be resized would have contradicted it.

29. The browser processor's version string is deliberately different. *Area: vision.* It is

browser-canvas-png-v1, not sharp-png-v1. Canvas and sharp do not resample or encode PNG identically, so two runs prepared differently must not be comparable by accident. Evidence records the preprocessing version for exactly this reason.

30. Arbitrary labels and licences. *Area: vision.* The metrics schema pinned perClass to `.length(VISION_LABELS.length) - seven`. Any dataset with a different number of classes ran to completion and *then* failed validation. Six fields were widened in one sweep, TypeScript interfaces and zod schemas together, after three separate build round-trips caused by widening one without the other. All three pre-existing evidence files were re-parsed afterwards to confirm backward compatibility.

31. The prompt has to come from the dataset. *Area: vision.* A deliberate failure run proved it: with the built-in prompt and a custom two-class dataset the result was 0% accuracy and 100% invalid output. `buildVisionPrompt(labels)` builds the closed label set from the dataset. Re-run with `red_panel/green_panel`: 100%, passed.

32. A loose file at the chosen root took the root folder's name as its class. *Area: vision.* `labelOf` now requires three path parts, so only `root/class/file` is accepted.

33. Four declarations block the run. *Area: vision.* SPDX licence text plus three privacy checkboxes. A dataset with no stated licence and no confirmation of what is in the images is not a dataset anyone should publish a metric from.

34. The real llava measurement, recorded as a failure. *Area: vision.* 21 samples, 14 correct: **66.7% accuracy, macro-F1 0.667**, 0% invalid output, 100% request success, median 529.6 ms, p95 721.4 ms. It **fails** its own gates (0.80 accuracy, 0.75 macro-F1) and the artefact records `"passed": false`. The per-class matrix closes exactly: `hardhat` has three false negatives against `safety_cone`'s three false positives - all three `hardhats` were called `safety_cones`, and `hardhat` scores F1 0.0. Two gloves and two masks were called `hardhats`. This is kept in the repository deliberately; see change 40.

3.5 Build, evidence and repository hygiene

35. .next held 142 stale files pointing at an old Downloads path against 75 correct ones, which is what the React Client Manifest errors were. Cleared.

36. .dockerignore was excluding something the build needed. *Area: DevOps.* Fixed and proven: `docker build -t edgepilot-ai:test .` completes in 154.8 s with `RUN npm run build` succeeding inside the image.

37. The evidence page traced the whole project into the standalone bundle. *Area: DevOps.* A dynamic `path.join(process.cwd(), root)` inside a loop defeats Next's file tracing. Rewritten with literal path segments.

38. /evidence exists as a page. *Area: DevOps.* `src/app/evidence/page.tsx` (286 lines, new) leads every artefact with `what_this_proves` and `what_this_does_not_prove` before showing a single number.

39. The ten-case evaluation matrix was missing. *Area: benchmark layer.* `scripts/evaluation/ten-case-matrix.ts` (340 lines, new), run with `npm run eval:matrix`, drives real module code - the same zod schemas, normaliser and readiness calculator the application uses - across 10 cases in 6 categories (1 normal, 3 malformed, 1 ambiguous, 3 injection, 1 missing-evidence, 1 provider-failure) and exits non-zero on failure. Current artefact: all 10 behaved as documented, no failures. Case 9 exists to assert that a missing hardware fit stays `null` and is never rendered as 0.

40. Working backups and the one real measurement were both mishandled by git. *Area: DevOps.* Nineteen `.bak` files sat beside real source, and ten of them were already tracked - so a `.gitignore` rule alone would not have removed them. Eighteen were cleared out and `*.bak` was added to `.gitignore`. In the same rule set, `/evidence/vision-benchmark/live-*.json` was silently ignoring the llava run - the first real measurement the project produced. Both decisions are now encoded as negation rules with the reasoning written beside them, rather than as `git add -f` flags someone has to remember.

41. scripts/db/ guards the shared database. Area: DevOps. `dotenv -e .env.neon -- prisma migrate deploy` fails *silently* when the file is absent: `dotenv` loads nothing, Prisma falls back to `.env`, and a command named "neon" runs against localhost while reporting success. `require-env-file.mjs` refuses to proceed unless the file exists, parses, and names a host that is not local; `neon.mjs` reads the file itself and prints the destination host before doing anything.

4. This structure against the team's

Compared with Desktop/Air on Accountants/edgepilot-ai-main (1)/edgepilot-ai-main, excluding `node_modules`, `.next`, `.git`, `package-lock.json` and `.env`. Fifty-five files differ, fifty-two paths exist only here - two of them local scratch (`.env.neon`, `_to_delete/`) rather than project content - and twenty-six exist only there.

4.1 What is gone

Removed	Why
<code>src/modules/device/</code> , <code>src/app/api/v1/devices/</code> , <code>src/shared/types/device.ts</code> , <code>tests/device/</code>	Self-reported hardware that no measurement consumed - see change 22
<code>src/app/auth/</code> , <code>src/app/api/auth/</code>	There are no accounts; ownership is a browser-generated session id
<code>sentry.client.config.ts</code> , <code>sentry.edge.config.ts</code> , <code>sentry.server.config.ts</code>	Three config files for a service with no DSN and no project behind it
<code>src/core/rate-limit/</code> , <code>src/core/security/</code>	A <code>rate_limits</code> table and a middleware layer guarding endpoints that have no authentication and no public URL
<code>src/data/</code>	A parallel data-access path beside <code>src/modules/*/infrastructure/repositories/</code>
<code>src/components/ErrorBoundary.tsx</code> , <code>Loading.tsx</code> , <code>LoadingStates.tsx</code> , <code>providers/</code>	Three overlapping loading components; the App Router's own <code>loading.tsx</code> and <code>error.tsx</code> conventions cover it, and <code>StateViews.tsx</code> holds the rest
<code>src/core/logging/logger.ts</code>	Superseded by <code>SessionLog.ts</code> , which has a redaction rule and an export format
<code>AUDIT_REPORT.md</code> , <code>FINAL_AUDIT_REPORT.md</code> , <code>PRODUCTION_AUDIT_REPORT.md</code> , <code>REMAINING_PRODUCTION_GAPS.md</code> , <code>TEST_DOCUMENTATION.md</code>	Five root-level audit documents describing the same codebase at five moments, none marked superseded. Replaced by this one record plus per-folder documentation
<code>public/index.html</code> , <code>public/favicon.svg</code>	A static page shadowing an App Router route
<code>docs/deployment.md</code>	Described a deployment that does not exist yet; the honest version is section 6

4.2 What is new

Added	What it does
<code>src/modules/benchmark/core/services/</code> - 8 new files, 1,761 lines	Hardware fit, modality, task compatibility, privacy class, egress warning, policy catalogue, comparison planner, comparison report
<code>src/modules/benchmark/infrastructure/OllamaCatalog.ts</code> , <code>OllamaResidencyProbe.ts</code>	Runtime status, installed models, and the memory reading that hardware fit is computed from

src/app/api/v1/comparisons/, local-runtime/, session-log/, session-log/share/	Four endpoints the original did not have
src/app/evidence/page.tsx	Artefacts readable without running anything
src/core/logging/SessionLog.ts, sessionLogStore.ts	Redacting, opt-in, exportable session record
src/lib/sessionOwner.ts, src/components/dashboard/session.ts	Account-free row ownership
src/components/dashboard/InstalledModels.tsx, VisionHandoff.tsx	Always-visible model panel with per-step filtering; handoff into the vision dashboard
src/components/vision/DatasetUpload.tsx and the two browser vision adapters	Bring-your-own dataset that never leaves the machine
scripts/db/, scripts/dev/, scripts/evaluation/	Shared-database guards, local helpers, the ten-case matrix
tests/core/, plus 7 new benchmark test files	279 tests across 25 suites, from 259
4 Prisma migrations	Nullable hardware fit and shared findings; optional email and session owner; drop devices; warm-up iteration
AI_USAGE.md	Required submission item, one section per work package

4.3 What changed in place

The twenty largest, by lines differing:

File	Theirs	Ours	Lines differing	Work package
src/app/globals.css	418	27	431	Dashboard
src/app/dashboard/page.tsx	692	15	697	Dashboard
src/components/dashboard/SetupPanel.tsx	322	201	379	Dashboard
src/app/vision-benchmark/page.tsx	73	244	293	Vision
src/modules/benchmark/application/services/BenchmarkRunner.ts	469	662	269	Benchmark
prisma/seed.ts	127	56	137	Integration
src/app/api/v1/workloads/route.ts	53	125	126	Integration
src/app/page.tsx	97	36	123	Dashboard
src/components/dashboard/RunPanel.tsx	214	315	119	Dashboard
src/components/dashboard/ProviderPanel.tsx	219	294	113	Dashboard
src/modules/benchmark/core/services/ReadinessCalculator.ts	40	88	110	Benchmark
src/components/dashboard/DashboardApp.tsx	134	203	91	Dashboard
src/modules/benchmark/application/use-cases/RunBenchmark.ts	316	308	90	Benchmark
tests/vision-benchmark/vision-execution.test.ts	301	361	86	Vision
prisma/schema.prisma	167	199	82	Integration
src/components/dashboard/api.ts	202	242	78	Dashboard
tests/benchmark/run-benchmark.test.ts	488	483	73	Benchmark
src/components/dashboard/RunResults.tsx	358	413	69	Dashboard
src/modules/vision-benchmark/application/executor.ts	159	222	65	Vision
src/modules/vision-benchmark/core/types.ts	159	201	62	Vision

`src/app/dashboard/page.tsx` going from 692 lines to 15 is not deletion. The route file is now a route file - metadata and a mount point - and the 692 lines of behaviour live in `src/components/dashboard/`, where they can be tested without a request.

5. Why this structure is the better one

5.1 The boundaries are real, not decorative

Both versions describe themselves as hexagonal. In this one the description holds when you check it. `src/modules/benchmark/core/` imports nothing from `infrastructure/`; the ports are two interface files, `AIProvider.ts` and `BenchmarkRepository.ts`; and `infrastructure/container.ts` is the single place where an adapter is constructed. No API route in this tree builds a provider, reads an environment variable, or touches Prisma. That is what makes the provider layer testable without a network: `tests/benchmark/helpers.ts` injects a fake `fetch` and a stepped clock, and the suite is deterministic because nothing underneath it can reach out.

The version with `src/data/` beside `src/modules/*/infrastructure/repositories/` had two answers to "where does data access live", which is the same as having none.

5.2 Every number carries its provenance

measured, derived, declared, unverified are values in the type system, not comments. A reader can tell a stopwatch reading from an arithmetic derivation from a placeholder without leaving the JSON. The rule that follows from it is the one that costs the most to keep: **null is never coerced to 0**. Reporting 0 ms for "we never got an answer" is the single most misleading number this system could produce, and there is a test file (`tests/dashboard/format.test.ts`) whose entire job is to stop it happening in the UI, and a case in the ten-case matrix that asserts it in the data.

A structure where `hardwareFit` was the constant 50 could not have supported that rule, because the field looked like a measurement and was not one.

5.3 It deletes

The clearest structural difference is that one tree accumulates and the other removes. Five audit reports at the root, three Sentry configs with no DSN, a rate limiter in front of unauthenticated endpoints, three loading components - each of those was reasonable when it was written, and none of them was ever taken back out. Dead structure is not free: it is read by reviewers as current, it appears in searches, and it makes the live parts harder to find.

Removing the device module cost a migration, a UI change, a module deletion and a test deletion. Leaving the column nullable would have cost nothing today and misled every reader afterwards.

5.4 The tests aim at refusals

The interesting question for a measurement tool is not "does it work" but "what does it refuse to say". So the suite is written against refusals: refuse to call a provider for a workload that does not exist; refuse to attribute a run to the wrong user; refuse to discard a measurement because the database went down mid-run; refuse to declare a comparison winner inside the noise; refuse to report a mean when the only iteration was a cold start.

That framing is also why `npm run build` is a required check alongside `npm test`. Jest type-checks only what it imports. `next build` type-checks the whole project, and it is what caught the two dev scripts in change 24 and the six schema-versus-interface mismatches in change 30 - nine defects that a green test suite could not see.

5.5 Evidence is an artefact, not a claim

`evidence/` holds files a reviewer can read without installing anything, and each leads with what it does **not** prove.

The vision run that fails its own gate is committed *because* it fails: a benchmarking tool whose only published measurement is a pass is not demonstrating a benchmark, it is demonstrating selection.

5.6 The client/server line is load-bearing

infrastructure/config.ts refuses to evaluate in a browser bundle, because every value it holds is a secret or a host. The two browser vision adapters exist for the inverse reason: images must *not* travel to a server, so the processor and the provider run in the tab. Both directions come from the same rule - decide where a thing is allowed to be, then make the wrong location fail loudly.

6. What this still does not do

Listed so nobody has to discover it in a review.

- **No cloud provider path has ever executed.** One Gemini or Groq key would light up the egress warning, tier-aware privacy and free-versus-paid comparison in a single run.
- **The privacy catalogue is still null and unverified.** Deliberate - see change 20 - and a human task.
- **official_source never reaches a recommendation**, which leaves an acceptance criterion open.
- **quality_score exists nowhere**, and is a required structured field.
- **No public URL.** The Docker image builds; nothing is hosted.
- **No UI** for comparisons, session-log or session-log/share, though all three endpoints work.
- **No adversarial or injection tests in the Jest suite.** The ten-case matrix covers the artefact side; the unit suite does not.
- **Two deliberate conflicts with the specification**, recorded rather than hidden: the device profile registry was a mandatory MVP item and was removed, and readiness_score() was specified to weigh privacy, which it no longer does. Sections 3.1 and 3.3 are the defence of both.
- docs/internal/database.md **is stale** - line 147 still says users own devices, and it omits shared_findings, session_id, warmup and privacy_class.
- device-profile.json **is still at the root**, referenced only by .dockerignore and two historical evidence files.
- describeEgressWarning **is tested in two files.** Flagged and left alone deliberately.

7. Verification state at the time of writing

Check	Command	Result
Lint	npm run lint	0 problems
Tests	npm test	279 passed, 25 suites
Build	npm run build	clean, 14 routes
Container	docker build -t edgepilot-ai:test .	succeeded, 154.8 s
Evaluation matrix	npm run eval:matrix	10/10 cases behaved as documented
Vision, live	vision:run:ollama --model=llava:latest	ran, 66.7%, failed its gate , recorded

Two credentials for the shared Neon database were exposed during the work that produced this record and should be rotated. .env and .env.neon are ignored by .gitignore line 28 (.env*); git status --short must show nothing matching .env before any commit.

8. Where the rest of the documentation is

Every folder in this project carries a FOLDER.md describing what it is, what happened in it, what each file inside it

does, and which other folders it is connected to. Those cross-references are written in both directions, so a trail can be followed either way.

Start at `src/modules/benchmark/FOLDER.md` for the measurement core, or `src/app/FOLDER.md` for the request path.

Seventy-five folder documents, 478 cross-folder links, every one of them resolving in both directions - if folder A points at folder B, B carries a "Referenced from" entry pointing back at A.

A PDF set lives in `Documents/AoC Documents`: this record on its own, twelve volumes grouped by area, and one complete 112-page bundle containing everything.