

The Test Suite

279 tests across 25 suites, and what they refuse to let the system say

EdgePilot AI - Team 15
Documentation set, 23 August 2026
5 documents in this volume

tests/ - the Jest suite

Work package: shared, by subfolder, matching the module each suite covers.

What this folder is

279 tests across 25 suites, mirroring `src/` one level down. Nothing here touches a network, a real timer, a real model or a real database - every adapter takes its `fetch` and its clock by injection, which is what makes the suite deterministic. A benchmark test that depended on real latency would be a flaky test that also proved nothing.

What happened here

Eight files were added and six changed, taking the suite from 259 tests to 279. `tests/core/` is new; so are the suites for hardware fit, task compatibility, privacy assessment, comparison, comparison runs, the Ollama catalogue and the egress warning.

The framing is worth stating because it explains why the tests look the way they do. For a measurement tool the interesting question is not "does it work" but "**what does it refuse to say**". So most of these assert refusals: refuse to call a provider for a workload that does not exist, refuse to attribute a run to the wrong user, refuse to throw away a measurement because the database went down mid-run, refuse to declare a comparison winner inside the noise, refuse to turn a `null` into a `0`.

What this suite cannot see. Jest type-checks only what it imports. Two dev scripts and six schema-versus-interface mismatches were invisible to a green run and were caught by `next build`, which type-checks the whole project. That is why both are required checks.

One duplication is known and deliberate: `describeEgressWarning` is covered in both `egress-warning.test.ts` and `privacy-assessor.test.ts`. It is flagged rather than consolidated.

Subfolders

Subfolder	Covers
<code>benchmark/</code>	19 files - the measurement core, the adapters, the routes
<code>core/</code>	The session log
<code>dashboard/</code>	The display helpers that enforce null-honesty
<code>vision-benchmark/</code>	5 files - the vision module end to end

Connected folders

- `src/` - the tree this mirrors.
- `.github/workflows/` - runs `npm test` on every push.

Referenced from

These folders point here. Each link below resolves in both directions.

- `src/modules/`

tests/benchmark/ - the measurement core under test

Work package: benchmark and provider layer.

What this folder is

Nineteen files covering the provider adapters, the runner, the use cases, the core services and the API route.

What happened here

Seven files are new, all of them covering services that did not exist before: hardware fit, task compatibility, privacy assessment, comparison planning, comparison runs, the Ollama catalogue and the egress warning.

`run-benchmark.test.ts` changed most among the existing files (73 lines differing) - it uses the **real** `BenchmarkRunner` over scripted providers rather than a stubbed runner, so the envelope that comes out is the one the API actually returns, and the cold-start change moved through it.

One fix in here is worth calling out because of how it was made rather than what it was: a run of three iterations was reporting "4/4 iterations succeeded", because the success chain was counting the discarded warm-up. The fix went into the runner, not into the assertion.

Files

File	What it covers
<code>helpers.ts</code>	The test doubles: fake providers, a stepped clock, a frozen clock, SSE and NDJSON response builders, connection-refused and hanging fetch stubs
<code>benchmark-runner.test.ts</code>	When the runner falls back. The rule: fall back only when the provider failed for a reason another provider could plausibly not share
<code>run-benchmark.test.ts</code>	The use case - ownership refusals, persistence failure mid-run, and the full envelope
<code>benchmarks-route.test.ts</code>	Parse, dispatch, status code. Above all: refuse a malformed body before anything constructs a provider, because a provider call costs money, quota and up to five minutes
<code>benchmark-measurement.test.ts</code>	What the summary refuses to report. This is where a number would be invented if anywhere
<code>benchmark-config.test.ts</code>	Environment validation and the server-side guard
<code>hardware-fit.test.ts</code>	VRAM residency, spill share, and the states between
<code>task-compatibility.test.ts</code>	The asymmetric rule - vision can do text, text cannot do vision, embedding neither
<code>privacy-assessor.test.ts</code>	Privacy class, disqualifiers, egress classification
<code>egress-warning.test.ts</code>	The pre-flight warning, including that it does not carry the prompt text
<code>comparison.test.ts</code>	Planning and reporting: overlapping ranges must not produce a winner
<code>run-comparison.test.ts</code>	The comparison use case, including sequential execution for local entrants
<code>ollama-catalog.test.ts</code>	Runtime status and the installed-model list
<code>ollama-provider.test.ts</code>	NDJSON streaming, nanosecond conversion, TTFT
<code>gemini-provider.test.ts</code>	SSE parsing, <code>usageMetadata</code> , header-carried key

<code>groq-provider.test.ts</code>	OpenAI-compatible SSE, <code>stream_options.include_usage</code> , <code>completion_time</code> in seconds
<code>demo-provider.test.ts</code>	The one promise that makes the demo adapter acceptable: nothing it produces can be mistaken for a measurement
<code>provider-errors.test.ts</code>	The published error vocabulary. This is what stops the code and <code>docs/local-model-setup.md</code> drifting apart
<code>provider-registry.test.ts</code>	Which adapters are usable in a given environment, and in what order

Connected folders

- `src/modules/benchmark/` - the module under test.
- `src/app/api/v1/benchmarks/` - the route `benchmarks-route.test.ts` drives with the container mocked.
- `docs/local-model-setup.md` - the error-code table `provider-errors.test.ts` pins.

Referenced from

These folders point here. Each link below resolves in both directions.

- `src/app/api/v1/`
- `src/app/api/v1/comparisons/`
- `src/app/api/v1/local-runtime/`
- `src/modules/benchmark/application/`
- `src/modules/benchmark/application/dtos/`
- `src/modules/benchmark/application/services/`
- `src/modules/benchmark/application/use-cases/`
- `src/modules/benchmark/core/`
- `src/modules/benchmark/core/ports/`
- `src/modules/benchmark/core/services/`
- `src/modules/benchmark/infrastructure/`
- `src/modules/benchmark/infrastructure/providers/`
- `tests/`

tests/core/ - cross-cutting concerns under test

Work package: benchmark and provider layer.

What this folder is

One suite, covering the session log.

What happened here

New, alongside `src/core/logging/`. The tests exist mainly to pin the redaction rule, which is the part of the session log that would be easiest to relax by accident: prompt **content** is never recorded by default, only a length and a digest. A log that wrote out the same prompt the egress warning had just flagged would have defeated its own warning.

They also pin the opt-in behaviour: a request without a session header is not logged anonymously and not logged to a default bucket - it is not logged at all.

Files

File	What it covers
<code>session-log.test.ts</code>	224 lines. Redaction, prompt digests, the share payload and its consent statement, and the in-memory store's bounds

Connected folders

- `src/core/logging/` - the code under test.
- `src/app/api/v1/session-log/` - the endpoints that expose it.

Referenced from

These folders point here. Each link below resolves in both directions.

- `src/core/`
- `tests/`

tests/dashboard/ - display helpers under test

Work package: product UI.

What this folder is

One suite, and it defends one rule.

What happened here

The wire contract makes every aggregate nullable on purpose, because reporting 0 ms for "we never got an answer" is the single most misleading number this system could produce. These tests pin the helpers that stop a null becoming a 0 on screen - which is the last place the rule could be broken, after the data layer has got it right.

Files

File	What it covers
format.test.ts	77 lines. fmtMs, fmtNum, fmtPct, fmtElapsed, describeErrorCode, isUuid, isPlaceholderId - each asserted to render "not measured" rather than a zero

Connected folders

- src/components/dashboard/ - format.ts, the file under test.
- evidence/evaluation/ - case 9 of the ten-case matrix asserts the same rule one layer down, in the data.

Referenced from

These folders point here. Each link below resolves in both directions.

- src/components/
- tests/

tests/vision-benchmark/ - the vision module under test

Work package: vision benchmark.

What this folder is

Five suites covering the vision workload from schema to provider.

What happened here

`vision-execution.test.ts` changed most (86 lines differing), because the warm-up classification changed the shape of an execution.

That change is a useful record of how a small feature goes wrong. The warm-up consumed a deterministic test clock, so the timings moved - fixed by calling `provider.classify()` directly rather than the timed wrapper. It fired on cloud providers, which meant a billable call and one extra image sent to a third party - fixed by restricting it to local providers and adding an explicit `warmUp?: boolean`. And it ate the first of the scripted `outputs.shift()` responses in two existing tests - those two now opt out, with a comment saying why, and two new tests cover the warm-up behaviour properly.

`vision-api.test.ts` changed to accept datasets whose label set is not the built-in seven.

Files

File	What it covers
<code>vision-benchmark.test.ts</code>	356 lines. The workload end to end: metrics, thresholds, evidence shape
<code>vision-execution.test.ts</code>	362 lines. The executor, including warm-up behaviour and its opt-out
<code>vision-providers.test.ts</code>	238 lines. The Ollama and Gemini vision adapters
<code>vision-infrastructure.test.ts</code>	204 lines. Image processing, manifest loading, the evidence store
<code>vision-api.test.ts</code>	78 lines. The <code>/api/v1/vision-benchmarks</code> route

Connected folders

- `src/modules/vision-benchmark/` - the module under test.
- `datasets/vision-benchmark/` - the fixtures pretest regenerates before these run.
- `src/app/api/v1/vision-benchmarks/` - the route.

Referenced from

These folders point here. Each link below resolves in both directions.

- `src/modules/vision-benchmark/application/`
- `src/modules/vision-benchmark/infrastructure/`
- `tests/`